

In a single pass, most local 4-bit executors tested transcribe a detailed plan; two of thirteen bind it

Abstract

We tested plan-conditioned execution of a mitochondrial variant-calling workflow. Twelve local 4-bit open-weight models and three commercial API models generated shell scripts in a single pass. Without a plan, 6 of 108 local runs and 13 of 27 API runs produced the correct call set. The API models succeeded with a short plan that contained no commands. The local models reached 34 of 36 correct runs only with a detailed plan that supplied literal commands. The model groups also differed in quantisation, size, vendor, serving system, and decoding parameters.

Runnable command syntax was associated with higher success than a non-runnable form of the same command. Removing explanatory prose produced no detectable difference. A moved-reference-path test then separated transcription from path binding. Eleven of thirteen local models copied the stale path and failed. Two used the path stated in the input description. When the execution error was returned, three of the eleven copying models repaired the path in at least two of three generations. Eight did not.

An RNA-seq case study produced benchmark-matching counts. It required human input preparation and did not complete the reporting step. The results apply to workflows whose plans specify all analytical choices.

Introduction

Routine sequencing analyses require an operator to connect a fixed workflow to new files, execute it, detect failures, and repeat failed steps. This study measured whether language models can perform these execution tasks from a separately authored plan.

The study used two workflows. Benchmark 1 used per-sample mitochondrial variant calling from paired-end Illumina reads. Benchmark 2 used per-sample RNA-seq quantification on a public Galaxy server. The study tested fixed model sets on five machines. It did not test analysis design or measure human time. Neither benchmark achieved unsupervised operation.

Stable pipelines do not require a language model. A validated shell script, workflow rule set, CWL file, or saved Galaxy workflow is a better choice (Di Tommaso et al. 2017; Mölder et al. 2021). These tools are deterministic and auditable. They also do not require a GPU. A language model can be useful for a one-time analysis or a change in file layout or tool version. It can also give a natural-language interface to a user who does not write code. Another possible use is *binding*. Binding connects a fixed workflow to a specific set of local inputs. Benchmark 2 tested some parts of binding (Results, *Case study*). The perturbed-plan test examined path binding directly. Two of thirteen local models changed the path correctly (Results, *A one-path perturbation*).

Workflow systems separate a specification from its execution. An engine connects the specification to concrete inputs (Di Tommaso et al. 2017; Mölder et al. 2021; Amstutz et al. 2016; Galaxy Community 2024). Related language-model methods include plan-and-solve prompting, least-to-most decomposition, ReAct, HuggingGPT, graph dispatch, and Reflexion (Wang et al. 2023; Zhou et al. 2023; Yao et al. 2023; Shen et al. 2023; Kim et al. 2024; Shinn et al. 2023). Alam and Roy (2025) reviewed applications to bioinformatics workflows. This study measured this design on genomic workflows. Plan detail was the independent variable in benchmark 1.

In practice, the execution side is supplied by an *agent harness*. Several were in wide use by mid-2026, and no comparison among them was run here. Benchmark 2 used `galaxyproject/loom`, for deployment reasons given in the Supplement. Nothing in the results is known to be harness-independent. Benchmark 1's harness is not an agent harness at all. It is one prompt in, one script out, executed externally, with no loop, no observation and no repair. That isolates plan detail as the only variable, but it means the system should not be called agentic. Benchmark 2 was added to supply the missing loop. The closest bioinformatics work evaluates open-ended tool-using research agents against task suites (Huang et al. 2025; Mitchener et al. 2025; Su et al. 2025; Mehandru et al. 2025). We ran none of those suites, so **no external anchor was obtained**

and no number here sits on a published scale. Benchmark 2 has two external comparators, both used with their protocols stated. First, a human executed the same workflow manually on the same server. Second, eight frontier models acting as *both* planner and executor on the same data cost \$2.82–\$131.83 per run (Nekrutenko 2026). That source is a post by the present author, with unpublished plans and no replicates. It bounds API cost and is not a frontier control.

Two properties of open-weight models help determine whether one fits on a given machine, and parameter count alone captures neither. A mixture-of-experts (MoE) model routes each token to a subset of experts. Its per-token compute is governed primarily by the active parameters, whereas weight residency is governed by the total parameters. Key-value-cache memory grows linearly with context length and also depends on architectural quantities including layer count, key-value-head count and head dimension. Grouped- and multi-query attention reduce this cost relative to full multi-head attention (Table 1).

Table 1. The four local models on which residency was measured, plus the API executors. This is a selection, not a catalogue. Fourteen local models appear in this study. The full list, with digests, declared context lengths and per-model default decoding parameters, is Supplementary Table S8. Local models were served by Ollama at 4-bit quantisation. Quantisation is not uniform: `gpt-oss:20b` ships MXFP4, while every other local model is `q4_K_M`. **No residency column is given**, because this study’s own residency figures for one of these models do not reconcile. Every residency measurement made, with the discrepancy stated inline, is Supplementary Table S4. Residency was measured for four of the fourteen models, and not for the other ten. Filling that column costs roughly six minutes of cold load per model. That was not done, and no figure is estimated. Declared parameter counts span 4.0 B (`nemotron-3-nano:4b`) to 36.0 B (`qwen3.6:35b-a3b`).

Executor	Architecture	Total params	Active params/token	Quantisation	Role
<code>qwen3.6:27b</code>	dense	27 B	27 B	<code>q4_K_M</code>	benchmark 1; cross-platform arm; error matrix
<code>qwen3.8:27b</code>	dense	27 B	27 B	<code>q4_K_M</code>	out-of-sample addition; benchmark 2
<code>qwen3.6:35b-a3b</code>	sparse MoE	35 B	3 B	<code>q4_K_M</code>	dense arm benchmark 1; benchmark 2
<code>gemma4:12b</code>	dense	12 B	12 B	<code>q4_K_M</code>	MoE arm throughput measurement
<code>claude-opus-5</code> , <code>claude-sonnet-5</code> , <code>claude-haiku-4-5</code>	commercial API	—	—	—	only frontier reference gradient
<code>claude-opus-4-7</code> , <code>claude-sonnet-4-6</code> , <code>claude-haiku-4-5</code>	commercial API	—	—	—	error matrix

Two workflow classes were run: alignment and small-variant calling, at 7 measured plan steps, and RNA-seq quantification, at 30 measured workflow steps. Four further classes cover most of the remaining routine per-sample work: single-cell RNA-seq, de novo assembly, metagenomic profiling, and ChIP-seq or ATAC-seq peak calling. None of the four was run. Which two classes were run is the single most important limitation in this paper. **The two classes run here are the two in which the plan can state the answer, and the four not run are ones in which it cannot.** In each of the four, at least one parameter is a scientific judgement made against the data: filtering thresholds, clustering resolution, assembler parameterisation, database version, peak-calling thresholds. For those classes, there is no canonical value for a plan to carry, and no invocation for an executor to transcribe. The plan-sufficiency result is therefore conditional on tasks with no data-dependent parameter choices. It may be in part a restatement of the task-selection criterion (Supplementary Table S19).

Two limits apply to all results. First, the plan gradient is exploratory. We wrote the most detailed plan after we observed failures with the shorter plans. Second, benchmark 2 is a case study. We do not combine results

from the two benchmarks. The detailed plan also creates a design confound. Its meta-prompt requests every flag and its exact value. It states that the executor will copy the commands into a shell script. Thus, success with v2 can include direct transcription. We measured transcription in two ways. A token-recall index gives an indirect measure (Results, *How much of the plan is copied*). A moved-path test gives a direct measure (Results, *A one-path perturbation separates transcription from binding*).

Methods

Extended methods are in the Supplement.

Study design

The benchmarks used different harnesses, models, scoring methods, and infrastructure (Figure 1). In both benchmarks, a commercial API model authored the plan before execution. A local 4-bit open-weight model executed the plan. Benchmark 1 used one prompt and produced one shell script. The model could not observe execution or retry. Benchmark 2 used multiple tool calls. The executor read state from a public server and kept state between sessions.




Figure 1: Study design

Figure 1. Study design. A plan is authored once, off the executor machine, by a frontier model in the planner role. Benchmark 1 executes that plan in a single pass: one prompt in, one bash script out, executed and scored with no observation step. Benchmark 2 executes it in a tool-using loop against the public usegalaxy.org server, with observation, durable state and repair. The dotted line marks the data-governance boundary as it was used here. The harness does not enforce that boundary. The frontier executor arms of benchmark 1 cross it, because their prompts carry the dataset manifest.

The initial benchmark-1 design contained 324 local generations and 81 API generations. It used 12 local executors, 3 API executors, 9 plan-file and track combinations, and 3 seeds. The nine combinations represent seven plan conditions because Track B occurs three times. Ten released logs contain 848 cells and 75.27 h of generation. API calls account for 81 cells and 1.55 h (Supplementary Tables S1–S2). Three later passes add 351 local cells. These are the 273-cell seed extension, the 39-cell perturbed-plan test, and the 39-cell repair test. The repair log has one duplicate skipped row, which is not counted as a cell. These later passes are not

part of the 848-cell ledger. We discarded 400 earlier runs that used a `q8_0` key-value cache. Supplementary Table S10 gives generation outcomes for each block.

Benchmark 1: task, plan conditions, and execution

The task was per-sample mitochondrial variant calling on four deeply downsized paired-end Illumina amplicon samples. The samples come from blood and cheek tissue of a mother-child pair, with nine truth variants per run. The reference workflow indexes the 16,569-bp chrM reference with `bwa index` and `samtools faidx`. It aligns with `bwa mem` piped to `samtools sort`, and indexes the BAM. It calls variants with `lofreq call-parallel`, then compresses and indexes the VCF with `bgzip` and `tabix`. It finishes with a `bcftools query` plus `awk` fan-in that collapses all four samples into one `collapsed.tsv`. Only the fan-in couples samples to one another (Supplementary Figure S1).

Executors were run under eight conditions, spanning no plan at all to a near-executable recipe (Table 2). Six conditions are plan files; Track B and v0.5 are prompt-template variants that substitute no plan file. Three planner meta-prompts produced the plan files. The meta-prompts are reproduced verbatim in the Supplement, with the executor prompt template. Two pieces of content that the results depend on originated in the human meta-prompt, not with the planner model. The first is the literal read-group string, which `PLANNER_PROMPT.md` requires reproduced “EXACT[ly], character-for-character”; it is therefore present in the *lean* v1 condition. The second is the naming of `lofreq call-parallel`, `bgzip` plus `tabix` and `bcftools query`. **The gradient is therefore in part a gradient in how much literal syntax was demanded of the planner, not purely in model-authored detail.** The exception to model authorship is v1g. Its `lofreq` block was extracted mechanically from the Galaxy IUC tool registry (tools-iuc commit 39e7456). It is a check on sensitivity to plan authorship, not an arm of a planner comparison.

Table 2. Plan conditions. Sizes are bytes of the substituted plan text, or of the prompt-template variant for the two no-plan conditions.

Condition	Size	Summary	Hypothesis tested
Track B	1.2 KB	Problem statement and tool inventory only	How much does a model do unaided?
v0.5	1.4 KB	Track B plus one line giving the tool order	Does sequencing alone help?
v1 (lean)	3.1 KB	Numbered bullets naming tools and key flags; no code	Reference lean plan
v1.25	3.1 KB	v1 plus one literal <code>lofreq call-parallel</code> command line	Does one full command bridge v1 to v2?
v1.5	1.3 KB	v2 with all prose stripped; code fences only	Does the prose change outcomes, or is it decoration?
v1g	4.2 KB	v1 plus a <code>lofreq</code> invocation extracted mechanically from the Galaxy IUC tool registry	Can a tool registry substitute for a plan author?
v2 (detailed)	4.6 KB	Exact command line and a Gotchas block per step	Reference detailed plan
v2_defensive	6.5 KB	v2 plus <code>try()</code> helper, per-step validation, retry-once, failure log	Does error-handling prose change runtime behaviour?

Every Track-B run is a no-plan run regardless of which plan file was nominally passed on the command line. Pooling by plan file without honouring this yields the convincing but false conclusion that the most detailed plan performs worst. `revision/scripts/gradient.py` implements the correct mapping and asserts it.

All local generations were served by Ollama 0.32.5 at `temperature = 0.2` with `num_ctx = num_predict = 16384`. Neither the local seed nor the API run identifier gives bitwise determinism. All claims of

reproducible sampling have been removed. **A confound, not merely a disclosure:** the harness overrode only `temperature`. It left `top_k` and `top_p` at each model’s own defaults. Those defaults are not uniform across families, and are absent altogether for `qwen3.8:27b` (Supplementary Table S8). Every cross-model comparison here therefore compares models decoded from different distributions. The gradient itself is a within-model comparison and is unaffected.

Generated scripts were executed on the model host, in a per-cell working directory, under a pinned conda environment, and killed at 600 s. Execution had no human review, and no container, VM, seccomp profile, namespace, cgroup quota or egress restriction. Only the wall-clock kill, the PATH-restricted inventory and the working directory are enforced. The constraints on absolute paths, package managers and network fetches are prompt instructions a model is free to violate. The violation rate is measured in the Results rather than asserted. **The local arm does not meet the sandboxing bar a production deployment needs and is not a deployment template.** Attacker- or accident-controlled text reached the executor through at least five unsanitised channels. No indirect-prompt-injection test was run.

Scoring, error injection, and interval estimates

Success was assessed on a three-rung ladder matching the released scorer’s field names. **M1** (`m1_executes`): `bash run.sh` returned 0 within 600 s. **M2** (`m2_schema`): every expected output exists and `bcftools` accepts every VCF. **M3** (`m3_jaccard`): the per-sample variant output matches the answer key. Two properties of M3 govern how the Results read. **M3 depends on exit status:** `score_run.py` sets it to 0.0 whenever the run exited non-zero, regardless of what is on disk. **M3 is the tolerant Jaccard overlap**, not a per-sample success fraction. It is the per-sample Jaccard of (chrom, pos, ref, alt) against the answer key, with allele frequencies matched within ± 0.02 , averaged over four samples. On the 324 local reasoning-off cells it takes only 0.0 and 1.0; on the frontier arm it does not. Conventional variant-calling metrics were computed alongside, not instead (Supplementary Table S12). The departures from the nearest published conventions (Krusche et al. 2019) are enumerated in Supplementary Table S5.

Tool failures were injected with PATH shims applied only to the targeted tool. The generated script saw nothing but ordinary exit codes, stderr and output files. Each executor–plan combination comprises 39 cells: 8 pattern–target combinations injecting a true error $\times$ 3 seeds, 4 injecting no error $\times$ 3 seeds, and 3 uninjected controls. **Two of the seven patterns do not inject an error at all.** `slow_tool` sleeps 30 s and then succeeds; `stderr_warning_storm` emits 200 warnings and then succeeds. A third of the injected matrix therefore tests latency and log-noise tolerance, not recovery. Results are reported by class and never pooled (Supplementary Table S6).

The full gradient was run at $n = 3$ seeds per cell. The headline condition (plan v2, Track A, reasoning off) was raised to $n = 10$ by adding seven seeds and pooling. After the second review cycle, the three discriminating conditions — v1g, v1.25 and v1.5, Track A, reasoning off — were likewise raised to ten seeds per model, across thirteen local executors. The thirteen are the twelve gradient models plus `qwen3.8:27b`. That model’s original three seeds sit in a separate log and are not pooled, so it contributes seven. One attempted v1g cell has no score record in the released log and could not be reconstructed; it is therefore missing rather than classified as a failed generation. The denominators are 126 for v1g and 127 for v1.25 and v1.5 (Supplementary Table S16). Treating the missing cell as a failure gives 31/127 for v1g and does not change the conclusions. This is **repeated sampling from a single configuration, not replication**. Ten seeds on one machine, one engine build, one quantisation, one plan and one scorer vary the sampler and nothing else. Generation attempts ending in truncation, refusal or provider error are scored 0 and retained; no completed generation was dropped for producing a bad answer. All success proportions carry Clopper–Pearson exact 95% confidence intervals. The intervals are wide here: 3/3 gives [0.29, 1.00], 10/10 gives [0.69, 1.00].

The perturbed-plan condition

One condition was added after the second review cycle to separate transcription from binding directly. The plan is v2 **verbatim** — not one character changed. The sandbox holds the same four read pairs, but the reference genome sits at `data/ref/GRCh38_chrM/rCRS.fa`. The plan’s literal path, `data/ref/chrM.fa`, does not exist. The prompt’s DATASET section states the real path, so every fact needed to succeed is in the

executor’s context. A script that copies the plan’s command lines fails at the first step. A script that binds the plan’s commands to the stated inputs succeeds. The sandbox was validated before any model ran, with a hand-written copier script (fails at `bwa index`) and a hand-written binder script (produces the correct call set). Thirteen local executors — the twelve gradient models plus `qwen3.8:27b` — ran at 3 seeds each, reasoning off, Track A. That is 39 cells (`matrix_jetson_perturbed.jsonl`; per-cell artifacts under `revision/runs_perturbed/`). The control is the unperturbed published v2 condition: 34 of 36 perfect across the twelve gradient models. No frontier executor was run on the perturbed plan, because no API spend was authorised for this pass.

Benchmark 2 and the fabrication study

The second benchmark re-quantified *Candidozyma auris* RNA-seq from Santana et al. (2023), BioProject PRJNA904261, six paired-end runs SRR22376027–SRR22376032. The target was a published benchmark with known expected values (~5,594 genes, ~92% assigned). Relative to benchmark 1, it varies four things at once: a second workflow class, a shared production server, a 30-step workflow, and an external cost comparator (Nekrutenko 2026). The executor-selection rule is invoked repeatedly, so it is stated verbatim:

take every local model that scored 3/3 at plan v2 Track A with reasoning off in the benchmark-1 gradient, and from that set take the largest dense model and the largest mixture-of-experts model by total parameter count.

The rule returns `qwen3.6:27b` and `qwen3.6:35b-a3b`. The dense arm actually run was `qwen3.8:27b`, substituted at matched parameter count after the rule was fixed. **That arm is therefore exploratory rather than confirmatory.** Plan sufficiency was never confirmatory on either benchmark, because the Galaxy plan was corrected repeatedly while benchmark 2 ran.

Both executors ran on one Jetson AGX Orin 64 GB, with reasoning suppressed, under `galaxyproject/loom`. Galaxy was reached over MCP, and `notebook.md` served as durable inter-session state. The plan was a router: a 50-line index plus seven step files, so the executor loads only the step it is on. The plan had seven steps, from importing the workflow by TRS identifier to reporting the per-sample counts. **No number reported for this benchmark comes from an agent’s own report.** Every state, parameter and count was read back from the Galaxy API or from dataset contents. The reason: one executor produced a fully formed failure report for an event that never occurred. Scoring had four components:

1. Invocation correctness, field by field.
2. Biological outcome against the published benchmark, read from dataset contents.
3. End-to-end reporting: did the executor itself deliver the per-sample table, with a dataset identifier beside every number?
4. A defect ledger attributing every human intervention to the planner or the executor side.

One change belongs in Methods because it changed the design. The original step-7 verification file stated its own expected values. That let an executor that never opens a dataset verify by recitation. The file was rewritten to state no expected values and to sanction `UNAVAILABLE`. The general rule: **any agent evaluation that places the expected answer in the task description cannot distinguish execution from recitation.**

The fabrication incident was a single observation, so a controlled follow-up bounded the rate. The design was 3 conditions $\times$ 2 models $\times$ 5 seeds, one fresh workspace per cell, a fixed prompt and known ground truth. The task was to fetch six featureCounts summary datasets from a real Galaxy history and report each **Assigned** value with the dataset identifier it was read from. The conditions were **blind** (nothing said about expected values), **leaky** (the expected range stated) and **blocked** (two identifiers replaced with same-shaped ones that do not exist, so `UNAVAILABLE` is the only correct answer). Answers were scored **CORRECT**, **WRONG** — the fabrication category — **ABSTAINED** or **ABSENT**. The design is **30 launches**. 24 launches produced a scorable answer file, giving 144 sample-level judgements; **six, 20%, produced no answer to score**. Three limitations govern the rate. Judgements within a session are not independent, so the session is the unit of analysis. Unscorable launches are excluded, so the rate is fabrication **conditional on an answer having**

been produced; the excluded class is structurally the one the original incident belongs to. And every cell is single-shot with a working tool path, whereas the incident arose in a resumed session with no reachable path.

Cost model, data governance, and availability

Two of the cost model’s eight inputs are measured. Frontier-API cost was \$0.0662 per run over 81 API cells. Jetson **wall-clock** time was 86.7 s per run; that includes model loading, execution and scoring, against 80.2 s of generation alone. The other six inputs are assumptions, exposed as parameters in `revision/scripts/cost_model.py` (Supplementary Table S9). The model reports a like-for-like **annualised** comparison, with an API-side labour term exposed rather than set to zero. Supervision labour is not priced at all, and benchmark 2 shows it is not zero. **The cost model is therefore the zero-touch case, which benchmark 2 did not achieve.**

The split also defines a data-governance boundary, and **it was respected by the planner role only**. The frontier *executor* arms — 315 API cells in all — transmitted the executor prompt. That prompt carries the dataset manifest with this study’s own subject labels. `revision/scripts/payload_audit.py` reconstructs every outbound payload and scans it against a deny-list of local-identity tokens (Supplementary Table S3). Benchmark 2’s MCP traffic is neither archived nor reconstructable, and it is outside the audit.

The repository at <https://github.com/nekirut/LLM-eval-paper> is released under an MIT license. It contains the plan files and prompt templates, the planner meta-prompts, the harness and sweep drivers, the scoring code, every archived `run.sh` and per-cell artifact, the error-matrix logs, and the revision analysis scripts. Four items are **not** released: the Galaxy plan files, the fabrication study’s code and per-cell outcomes, the eleven-item defect ledger with its diffs, and the `loom` session logs. Benchmark 2’s per-sample counts and gene count are re-readable from public Galaxy histories. Its defect count and attribution, its tool-call and wall-clock figures, and the fabrication study’s cell counts rest on `revision/galaxy_demo/FACTS.md`. That file is a hand-typed contemporaneous record of live API read-back — **the same evidentiary class as an assertion in this manuscript** — and those figures are marked as such wherever used.

Results

Results with plans that contain little detail

With no plan, 6 of 108 local runs produced the correct call set (exact 95% CI [0.02, 0.12]). The frontier executors produced the correct call set in 13 of 27 runs ([0.29, 0.68]). Plan v1 was the shortest plan file. It did not contain command lines. The frontier executors were correct in all 9 v1 runs. They were also correct in all 36 runs with the four more detailed plans. More plan detail did not increase their observed success proportion. This result has a wide interval. The exact 95% interval for 9/9 is [0.66, 1.00]. The local executors reached 34 of 36 correct runs only with v2 (Figure 2, Table 3).

The tested executor classes differed in how much specification they needed, with the largest contrast where the plan was thinnest. The local arm is 4-bit quantised, at most 35 B, open-weight and served by Ollama locally. The frontier arm is unquantised, of undisclosed size, from a single vendor, served over an API and decoded under different sampling parameters. These variables move together, and this design separates none of them. Quantisation in particular is a live alternative explanation for a model that cannot emit a correct `lofreq call-parallel` flagset. The decisive control — `qwen3.6:27b` at `q8_0` or `fp16` against its own `q4_K_M` — was not run. If precision closes a meaningful part of the gap, the finding is about quantisation rather than open weights. Three models from one vendor is also not a sample of frontier models.

Figure 2. Plan gradient with reasoning disabled, Jetson AGX Orin, $n = 3$ seeds per cell. Rows are the twelve benchmark-1 local executors plus the out-of-sample addition `qwen3.8:27b`, thirteen rows in all. The twelve-model gradient totals given in the text exclude `qwen3.8:27b`, which was run separately and is shown here for comparison. Columns run left to right from no plan (Track B) to the most detailed plan (v2). Cell values are the mean tolerant Jaccard score (M3), which on this arm takes only the values 0 and 1 per run.

Table 3. Plan gradient, reasoning off, twelve local executors, three seeds per cell. Track B pools three (plan file, track) cells, hence $n = 108$. Intervals are Clopper–Pearson exact 95% intervals. **The local arm is**




Figure 2: Plan gradient with reasoning disabled

reported as a count of runs producing the correct call set, not as a mean score. On all 324 local reasoning-off cells, M3 takes only the values 0.0 and 1.0. A local “mean M3” column would be the perfect-run proportion written twice. It would invite the reader to believe partial biological credit was measured, which on this arm it was not. The tolerant Jaccard is retained only for the frontier arm, where it is non-degenerate. Its cells there take the values 0.0, 0.833, 0.938 and 1.0, so a Track-B mean of 0.825 coexists with only 13 of 27 runs perfect. The frontier column pools the three API executors at $n = 9$ per column, 27 at Track B. Three local conditions — v1g, v1.25 and v1.5 — were later raised to ten seeds per model over thirteen executors (31/126, 63/127 and 109/127; Results, *Effects of prose and executable syntax*; Supplementary Table S16). The $n = 3$ rows are retained here so that the survey is one design.

Condition	Local runs with the correct call set	95% CI	Frontier runs with the correct call set	95% CI	Frontier mean M3 (Jaccard)
Track B (no plan)	6/108	[0.02, 0.12]	13/27	[0.29, 0.68]	0.825
v0.5 (tool order only)	1/36	[0.00, 0.15]	3/9	[0.08, 0.70]	0.801
v1 (lean)	9/36	[0.12, 0.42]	9/9	[0.66, 1.00]	1.000
v1g (registry-authorized)	9/36	[0.12, 0.42]	9/9	[0.66, 1.00]	1.000
v1.25 (v1 + one command line)	15/36	[0.26, 0.59]	9/9	[0.66, 1.00]	1.000
v1.5 (v2 with prose stripped)	30/36	[0.67, 0.94]	9/9	[0.66, 1.00]	1.000
v2 (detailed)	34/36	[0.81, 0.99]	9/9	[0.66, 1.00]	1.000

Effects of prose and executable syntax

v1.5 is v2 with all prose stripped, leaving only code fences. It is the **shortest** plan in the set, at 159 words against v2’s 660, and it is the second best.

The two halves of the prose-versus-syntax contrast are not equally well resolved. Removing the literal command lines, which is what v1 is, reduced perfect runs from 34/36 to 9/36. Stripping every word of explanation from v2 reduced them from 34/36 to 30/36, a difference **not resolved at these sample sizes**: paired by model, 3 improved, 1 worsened and 8 were unchanged (two-sided sign-test $p = 0.63$). The seed extension gave pooled descriptive proportions of 31/126 for v1g (0.246, exact 95% CI [0.17, 0.33]), 63/127 for v1.25 (0.496, [0.41, 0.59]) and 109/127 for v1.5 (0.858, [0.79, 0.91]). Because repeated seeds are clustered within models, the primary contrasts treat the 13 model-level success proportions as paired observations. The mean within-model difference was 0.253 for v1.25 minus v1g (exact sign-flip $p = 0.027$) and 0.354 for v1.5 minus v1.25 ($p = 0.0039$). The corresponding run-level paired counts are reported in Supplementary Table S16 as a sensitivity analysis. The defensible claim has three parts: runnable command syntax improved outcomes across the tested models; removing prose produced no detectable change; and the size of any prose effect remains uncertain.

The mechanism is more specific than a count of command lines. v1.25 and v1g each add one literal `lofreq` invocation to v1, but present it differently. v1.25 supplies one runnable line: `lofreq call-parallel --pp-threads 4 -f data/ref/chrM.fa -o results/{sample}.vcf results/{sample}.bam`. v1g supplies a six-line block in the registry’s long-form style, with two bare, valueless flags (`--sig`, `--bonf`) that the plan itself says may be omitted. Their pooled success proportions were 0.496 and 0.246, respectively. v1g also states that `results/{sample}.bam` is a positional argument at the end, so the contrast is not explained by omission of that information. Across the 13 paired model-level proportions, v1.25 exceeded v1g by a mean of 0.253 (exact sign-flip $p = 0.027$). This result is consistent with copy-pasteability, rather than information content or verbosity alone, as the relevant plan property. v1.5 is the shortest plan and the second best; v2 is the longest and the best; and v1g is the second longest but the weakest of the plans containing command syntax (Table 4). No rank correlation is quoted over five points with ties.

Table 4. Syntax density against outcome, Track A, reasoning off. “Literal command lines” counts non-blank lines inside fenced blocks that invoke a workflow tool or shell construct. That is an operational definition, and

the count is sensitive to it. v1g’s single `lofreq` invocation is written across six physical lines in the registry’s long-form style; it is counted as one, because it is one invocation. v1g’s author is the Galaxy IUC registry, extracted mechanically; every other plan was authored by a frontier model. Jetson values are recomputed under this manuscript’s retention rule, which scores the one truncated v1g generation 0 and keeps it. They therefore differ in the third decimal from `revision/galaxy_demo/SYNTAX_DENSITY.md`, which drops it. RTX 5080 values come from that file under its own convention. They are given only to show that the ordering reproduces on a second machine. Every frontier cell here is 9/9. The RTX 5080 sheet reports 0.667 in its frontier v1g cell, on a different set of 9 runs; neither figure should be over-read at $n = 9$. The ten-seed column pools the seed extension over thirteen local executors (Methods; Supplementary Table S16). v1 and v2 were not extended; v2’s headline pooling to $n = 10$ is Table S11.

Plan	Author	Words	Literal command lines	Local mean M3, Jetson ($n =$ 3/model)	Local mean M3, Jetson ($n =$ 10/model)	Local mean M3, RTX 5080	Frontier mean M3, Jetson
v1	frontier planner	422	0	0.250 ($n =$ 36)	—	0.211 ($n =$ 38)	1.000 ($n =$ 9)
v1g	Galaxy IUC registry (mechanical)	558	1	0.250 ($n =$ 36)	0.246 ($n =$ 126)	0.189 ($n =$ 37)	1.000 ($n =$ 9)
v1.25	frontier planner	413	1	0.417 ($n =$ 36)	0.496 ($n =$ 127)	0.528 ($n =$ 36)	1.000 ($n =$ 9)
v1.5	frontier planner	159	10	0.833 ($n =$ 36)	0.858 ($n =$ 127)	0.892 ($n =$ 37)	1.000 ($n =$ 9)
v2	frontier planner	660	8	0.944 ($n =$ 36)	—	0.949 ($n =$ 39)	1.000 ($n =$ 9)

Across the same five plans, the API executors were correct in 45 of 45 runs. The local success proportion increased from 0.250 to 0.944. Thus, greater plan detail reduced the observed difference between the two executor groups. The API result is a point estimate from nine runs per plan. v1g is not a planner-comparison arm because syntactic density was not matched. It is a mechanically extracted plan with a non-model author. Its result shows that authorship and plan length do not explain the observed pattern. The data support a narrower conclusion: local models produced working pipelines more often when the plan contained commands that could be copied without modification. Explanatory prose produced no detectable effect at the tested sample size. The moved-path test below measures direct transcription.

How much of the plan is copied

We define a **transcription index**: the token-level recall, in an emitted script, of the 90 literal command tokens of plan v2, after normalising the sample placeholder and the thread count. The corpus is the 787 reasoning-off scripts.

Table 5. Transcription index by plan condition, reasoning off, over the 787 archived reasoning-off scripts (the original corpus plus the seed-extension scripts at v1g, v1.25 and v1.5). Recall is of plan v2’s 90 literal command tokens, so all conditions are measured against one fixed reference. Brackets are the interquartile range across scripts. The corresponding bootstrap 95% intervals on the mean have a half-width of at most 0.03 in every cell with $n > 9$, and are omitted for width. The corpus includes the 48 `qwen3.8:27b` scripts that Table 3 excludes; they are 6% of it.

Condition	Local, all scripts	n	Frontier, all scripts	n	Local, correct runs only	n	Frontier, correct runs only	n
Track B (no plan)	0.446 [0.42–0.49]	117	0.520 [0.47–0.56]	27	0.428	6	0.505 [0.46–0.54]	13
v0.5	0.454 [0.42–0.49]	39	0.501	9	0.467	1	0.444	3
v1	0.515 [0.48–0.54]	39	0.522	9	0.509	9	0.522	9

Condition	Local, all scripts	n	Frontier, all scripts	n	Local, correct runs only	n	Frontier, correct runs only	n
v1g	0.520 [0.48–0.56]	129	0.517	9	0.505	34	0.517	9
v1.25	0.516 [0.49–0.54]	130	0.533	9	0.513	66	0.533	9
v1.5	0.702 [0.68–0.74]	130	0.658	9	0.705	112	0.658	9
v2	0.746 [0.74–0.76]	122	0.695	9	0.749 [0.74–0.76]	110	0.695	9

The 13 correct API runs with no plan had a mean transcription index of 0.505. Correct local runs with v2 had a mean index of 0.749. Thus, approximately half of the v2 command tokens occur in a correct pipeline without plan conditioning. Plan conditioning added approximately one quarter of the tokens. The index did not predict success within a plan condition. At v2, the means were 0.749 for successful scripts and 0.722 for failed scripts. At v1, the means were 0.509 and 0.517. Local scripts recovered 0.942 of the tokens in the code-only v1.5 plan ($n = 130$).

What the index cannot settle is what the remaining quarter is. A model that reflows supplied tokens and a model that binds a specification to inputs both score near 0.75. No token count forces them apart. The condition that does is a perturbed plan — v2 verbatim, with the reference moved — because a transcriber reproduces the defect and a binder repairs it. **That condition was run after the second review cycle, and it is the next section.**

A one-path perturbation separates transcription from binding

The design is in Methods. The plan is v2 verbatim. The reference genome sits at `data/ref/GRCh38_chrM/rCRS.fa` instead of the plan’s literal `data/ref/chrM.fa`. The real path is stated in the prompt’s DATASET section. The sandbox was validated with a hand-written copier (fails) and binder (succeeds) before any model ran. Thirteen local executors ran at three seeds each, against a control of 34/36 perfect on the unperturbed v2.

The result is 6 of 39 perfect, and the outcome is bimodal by model: every model is 3/3 or 0/3 (Table 6). This was verified from the emitted scripts, not inferred from the scores. Both surviving models wrote `data/ref/GRCh38_chrM/rCRS.fa` throughout. All eleven failing models wrote the plan’s `data/ref/chrM.fa` verbatim; several derived index names such as `chrM.fa.fa` from it. Every failing run died at the first step, with `bwa index` unable to open the file.

Table 6. The perturbed-plan pass: plan v2 verbatim, reference moved to a path stated in the prompt, thirteen local executors $\times$ 3 seeds, reasoning off. Outcomes read from the emitted scripts and per-cell logs (`matrix_jetson_perturbed.jsonl`; artifacts under `revision/runs_perturbed/`). Control: unperturbed v2, 34/36 perfect.

Outcome	Models	What every emitted script wrote
Bound (3/3 perfect)	<code>gemma4:31b</code> , <code>gpt-oss:20b</code>	the true path <code>data/ref/GRCh38_chrM/rCRS.fa</code> , carried into every derived path
Copied (0/3 perfect)	<code>qwen3.6:27b</code> , <code>qwen3.8:27b</code> , <code>qwen3.6:35b-a3b</code> , <code>gemma4:26b-a4b</code> , <code>glm-4.7-flash</code> , <code>granite4.1:30b</code> , <code>granite4.1:8b</code> , <code>laguna-xs-2.1</code> , <code>nemotron-3-nano:30b-a3b</code> , <code>nemotron-3-nano:4b</code> , <code>qwen3.5:4b</code>	the plan’s stale <code>data/ref/chrM.fa</code> verbatim; every run failed at the first step (<code>bwa index</code> : fail to open file)

Four things follow, and together they change what the gradient means. **First, for eleven of thirteen local models, v2 success is transcription.** They copy the plan’s command lines without reconciling them against the stated inputs. One path mismatch, stated in the prompt they were reading, sent a condition

that scores 34/36 to 0. Reviewer 2’s objection — that the headline result rewards pasting — is therefore confirmed for most of the model class, as a measurement rather than a concession. **Second, binding exists in this model class but is rare.** Two of thirteen models adjusted every path and produced the correct call set in all three seeds. Binding is therefore demonstrable, not hypothetical — and it is a property of individual models, not of the class. **Third, rank on the unperturbed gradient does not predict binding.** The paper’s original headline model, `qwen3.6:27b`, 10/10 at unperturbed v2, copied. `gpt-oss:20b`, one of only three models below 10/10 at the unperturbed headline condition (9/10), bound. **Fourth, the claim that survives is narrower than plan sufficiency: a detailed plan makes most local models work only while its literal paths match the data.** Reliable execution under a stale plan requires either a binding-capable model or a plan kept exactly current. Whether frontier executors bind under the same perturbation is unknown, because no frontier arm was run. And at three seeds per model, a 3/3 is [0.29, 1.00], so the bimodality is sharp but each per-model rate is wide.

Three attempts separate transcribers that repair from transcribers that do not

The perturbation above is single-pass, so a bounded repair arm re-ran it with feedback. The condition is the perturbed plan, unchanged. If the script exits nonzero, the model sees its own script, the exit code and the last 40 lines of the execution log, and may submit a fix, up to three attempts. The retry signal is the exit code only; the score is computed afterward and never shown. Thirteen local executors ran at three seeds each (`matrix_jetson_repair.jsonl`; per-attempt scripts under `revision/runs_repair/`).

Feedback split the eleven copiers into models that repair and models that do not (Table 7; per-seed attempts and scores are Supplementary Table S18). Three of the eleven copier models repaired. Eight did not, although the error named the missing file and the correct path was in their prompt. Repair rescued 8 of the 33 copier seeds. With repair, the class reached 14 of 39 perfect runs, against 6 of 39 one-shot and 34 of 36 unperturbed. The single-pass constraint therefore accounts for part of the transcription finding, not all of it. The taxonomy is three tiers: bind (2 models), repair (3), neither (8). One retry converted `qwen3.6:27b` and `qwen3.8:27b` from total failure to perfect. For those two models, an execute-and-retry round was worth more than any amount of plan prose.

Table 7. The bounded repair arm: the perturbed condition with up to three attempts, the error fed back after each nonzero exit. Thirteen local executors $\times$ 3 seeds, reasoning off. Per-model attempts and per-seed scores are Supplementary Table S18.

Tier	Models	Outcome
Bind on attempt 1	<code>gemma4:31b</code> , <code>gpt-oss:20b</code>	one attempt, perfect, all seeds (the controls; unchanged)
Repair on attempt 2	<code>qwen3.6:27b</code> , <code>qwen3.8:27b</code>	perfect on attempt 2 in all three seeds
Partial repair	<code>laguna-xs-2.1</code>	perfect on attempt 2 in two of three seeds
No repair	the remaining eight models	0 of 24 seeds recovered in three attempts

Three behaviours were verified from the per-attempt scripts, not inferred from the scores. First, `qwen3.6:27b`’s fix is a genuine repair. Its second attempt tests for the old path, falls back to the path stated in the prompt, and proceeds. It used the prompt’s information once the error pointed at the file. Second, `qwen3.6:35b-a3b` resubmitted the identical dead path in all three attempts, with the error naming that path in front of it. Third, `granite4.1:30b` and `gemma4:26b-a4b` guessed `data/ref/rCRS.fa` on attempt 3: the correct filename from the prompt, in the wrong directory. They used part of the prompt’s information and did not check the rest. Failure to repair is not failure to act. The non-repairers changed other things — flags, loops, index names — while keeping the dead path. Four caveats bound this arm: three seeds per model, one perturbation type, a retry only on a nonzero exit, and no frontier arm.

Repeated sampling with plan v2

The $n = 3$ headline does not survive repeated sampling. Seven seeds were added to the original three at v2 Track A (Figure 3, Supplementary Table S11). Nine of 12 executors are 10/10 (exact 95% CI [0.69, 1.00]). `gpt-oss:20b` is 9/10; `qwen3.5:4b`, which scored 3/3 at $n = 3$, is 7/10; `nemotron-3-nano:4b` is 1/10 ([0.00, 0.45]). Every model that moved, moved down. That is the expected direction rather than symmetric noise. We therefore report $n = 10$ as the headline and treat every $n = 3$ cell as an interval. The out-of-sample pair `qwen3.8:27b` against `qwen3.6:27b` is underpowered rather than null. 8 of 27 matched cells are discordant, splitting 6 to 2 in favour of the *older* model ($p = 0.29$). The generational framing has therefore been removed from the abstract and the conclusion.




Figure 3: Repeated sampling with plan v2

Figure 3. Repeated sampling with plan v2 (plan v2, Track A, reasoning off). Grey points are the $n = 3$ means as originally published; red points are the pooled $n = 10$ means. Every model that moved, moved down; no model rose. This is repeated sampling from one configuration, not independent replication.

The conventional metrics are not degenerate: recall without the exit-status rule sits above run-level success

Over all 324 reasoning-disabled runs, TP = 1,042, FP = 0 and FN = 1,874. Precision is 1.000 by construction, because the truth set is the canonical output of the same deterministic caller. **Recall, however, does not equal the run-level success rate** in any condition except v2. It is 0.086 against 6/108 at Track B, 0.306 against 9/36 at v1, 0.485 against 15/36 at v1.25, and 0.886 against 30/36 at v1.5 (Supplementary Table S12). Pooled, recall is 0.357 against a success rate of 0.321. No pooled F1 is reported: it averages over conditions that share no denominator of interest, and it re-expresses a binary.

The recall curve computed without the exit-status rule sits above the run-level success curve at every lean condition. The excess is real biology: variants correctly called by runs the run-level metric scores 0. Two mechanisms produce it. **26 of the 324 runs produced exactly 2 of the 9 truth variants**, concentrated at the lean end, so output is near-bimodal but not strictly so. And **8 of the 324 runs wrote all four VCFs and then exited non-zero**, of which 6 had recovered all nine truth variants. That is why 110 runs recovered the complete truth set while only 104 scored a perfect run. Scoring such a run 0 is defensible — a pipeline that exits non-zero has not finished — but it is a choice. That choice, not the biology, is what makes the run-level metric degenerate. So the conventional metrics quantify how often a model gets the biology right and the plumbing wrong: 8% of runs by partial call sets, plus 2.5% by exit status. What they cannot do is separate a good caller from a bad one. The truth set is the same caller’s own output, and no perturbation series that would vary accuracy continuously was run.

Enabling reasoning does not substitute for plan detail

Over 324 reasoning-disabled and 270 reasoning-enabled cells, restricted to the 10 models runnable in both arms, reasoning helps where the plan is thin and costs accuracy where it is sufficient. The per-plan differences are **+0.17 at v1**, **+0.23 at v1.25** and **−0.20 at v2**. At v2 Track A, 4 of 10 models regress and none improve. Enabling reasoning on `qwen3.8:27b` at plan `v1g` also needs a budget two to three times larger merely to emit a script. Even then, it falls from 3/3 correct in ~130 s to 1/3 correct in 7,395–12,974 s. **Chain of thought is not a substitute for a sufficient plan, and at the sufficient plan it is a liability.** The reasoning-enabled arm has a 6.7% harness-bounded failure rate against 0.3% disabled. It therefore measures model-plus-budget rather than model capability (Supplementary Results, Figures S2–S3).

Injected failures: the outcome is set by the plan and the pattern, not by the model

For `qwen3.6:27b` with plan `v2`, the 24 true-error cells produced 15 crashes, 6 propagated errors, and 3 recoveries (Figure 4). The recovery rate was 12.5% (exact 95% CI [0.03, 0.32]). The crash rate was 62.5% ([0.41, 0.81]), and mean M3 was 0.000. The other 12 cells did not contain a tool failure. Results that combine these two classes are not recovery rates. These single-pass runs could not retry. The repair experiment later tested retries with the moved-path condition (*Three attempts separate transcribers that repair from transcribers that do not*).

The identical distribution obtained for the frontier executor is a degeneracy, not a parity finding. Across the 39 v2 cells, `qwen3.6:27b` and `claude-opus-4-7` agree in **every cell on both handle category and M3**. The same holds under `v2_defensive`. Five of the six executors produce exactly crash 15 / propagate 6 / recover 3 on the 24 true-error v2 cells. Two architecturally unrelated families do not converge cell-for-cell by shared competence. The handle category at the baseline plan is therefore essentially determined by the injected pattern. Two things bound that reading. The per-cell artifacts for the error matrix were **not archived**, so the script diff that would confirm it cannot be run. And the frontier comparator here is a **different model generation** from the arm that produces the paper’s main result. The model term is not zero everywhere. Under `v2_defensive`, `qwen3.6:27b` and `claude-opus-4-7` resolve the 24 true-error cells to recover 9 and partial 15. `qwen3.6:35b-a3b` resolves them to recover 3 and partial 21, with its mean M3 *falling* to 0.208. `granite4` produces 24 crashes in 24 cells. **So a defensive plan changes error handling, and how much it helps is model-dependent.**

Figure 4. Injected-failure matrix: modal handle category across three seeds, for the local executor `qwen3.6:27b` and the frontier executor `claude-opus-4-7`. Top block: baseline plan `v2`. Bottom block:




Figure 4: Injected-failure matrix

plan v2_defensive. The two executors match in all 39 cells of each block. Five of the six executors in the matrix match on the v2 block. That is evidence that the outcome is pattern-determined, not evidence of model parity (see text). Columns 8–11 (`slow_tool@bwa`, `slow_tool@lofreq`, `stderr_warning_storm@bwa`, `stderr_warning_storm@lofreq`) inject no error; they are latency and log-noise tolerance checks, not recovery tests (Supplementary Table S6). The columns are alphabetical, so the twelfth and rightmost, `wrong_format_output@lofreq`, is a true-error pattern, and a green cell there is a genuine recovery.

A lexical audit of the 766 archived `run.sh` files replaces the claim that no destructive or escape behaviour was observed. None invokes a package manager, network fetch tool or `sudo`. None directs output to an absolute path outside `/dev/` and `/tmp/`. 17 (2.2%) write under `/tmp/`. One script — `rm -f /tmp/*.txt /tmp/temp_*.vcf.gz` — is a genuinely out-of-sandbox action, and it came from a no-plan run. The `/tmp/` writes concentrate in no-plan runs generally: 11 of 303 no-plan scripts, against 6 of 463 written with a plan. A model given no plan invents its own scratch-space convention, and the convention it invents leaves the sandbox.

Case study: a tool-using agent loop on a production server

This is a case study, not an experiment. It has one invocation per arm, no replicates, no frontier control, and a plan repeatedly corrected while it ran. Nothing in it carries a confidence interval. Tool-call and wall-clock figures are transcribed from `FACTS.md`; the counts are re-readable from the public histories.

Given the plan and a live MCP connection to usegalaxy.org, both local executors imported the 30-step IWC `rnaseq-pe` workflow by TRS identifier. Both verified the six paired-end inputs and the annotation. **A human had pre-built both of those after a harness transport fault, so those two plan steps were verification and not construction.** Both executors then submitted a parameter-correct invocation: 2 of 2 inputs and 8 of 8 parameters. The dense arm did this in 38 tool calls over 44 min, the MoE arm in 19 over 22 min. **The 8/8 figure measures transcription, not derivation**, because the invocation step supplied a literal JSON body. The invocations reproduced the published quantification. Both arms yielded 5,594 genes, verified by downloading and counting the counts table rather than by accepting a reported figure. Assigned fragments were 15.7–25.3 M per sample, at 93.1–93.6% of uniquely mapped fragments, against a published ~92% on that denominator (Supplementary Table S14, Figure 5). The two arms produced byte-identical values from twelve distinct dataset identifiers. That is expected from a deterministic server-side workflow on identical inputs, and it is not independent replication of model behaviour.

Figure 5. Per-sample featureCounts results for the benchmark-2 case study. Left: assigned fragments per run; the dense (qwen3.8:27b) and mixture-of-experts (qwen3.6:35b-a3b) arms are exactly superimposed.




Figure 5: Per-sample featureCounts results

Right: assignment rate against two denominators — uniquely mapped fragments and all fragments — with the published ~92% benchmark marked. The gap between the two denominators is multi-mapping, a property of the data rather than of the executor.

The result carries three caveats that must accompany it wherever it is claimed. Both executors submitted a parameter-correct invocation, and the workflow produced benchmark-matching counts. **But the plan supplied the invocation body literally, a human pre-built two of the seven plan steps’ inputs after a transport fault, and neither executor completed the final reporting step.** Neither reported a single per-sample value in the two runs of step 7. A human read the six values out of the server (*Task scoping*). “Both executors submitted a correct invocation” is true; “both executors succeeded” is not. Supervised human time was substantially larger than the agent wall time and was not logged. Reasoning was suppressed in both models. Job-level outcome counts are not reported, because the jobs-summary JSON was not archived.

Repeated human correction of the plan, and not any change of model, is what produced a clean run. FACTS.md records eleven plan defects, all authored on the planner side, against two executor errors and one harness error. **The ledger and the diffs that fixed each entry are not released.** The attribution therefore cannot be checked, and must be taken on the author’s contemporaneous record. That is why the count is used here and not in the Introduction, Discussion or Conclusion, which say “repeated human correction” instead. The ledger, and the five defect classes worth naming for plan authors, are Supplementary Table S15 and its accompanying text. The most expensive defect is the clearest instance: a GFF3 annotation staged where featureCounts requires a GTF carrying `gene_id`. It aborted all six counting jobs two hours in, because the IWC workflow hard-codes the `exon/gene_id` pair and exposes neither as an input. Plan sufficiency was established during this exercise rather than tested by it. The case study contains no plan gradient, so nothing in it discriminates a sufficient plan from an insufficient one.

Failure modes visible only in the agentic setting

Three failure modes appeared that a single-pass benchmark cannot produce. First, neither executor completed the final reporting step, so a scorer reading only the invocation would record two clean successes. Second, neither model could escape its own bad durable state. One arm re-read its notebook, found the annotation identifier recorded before the GFF3 was replaced, and re-invoked the doomed workflow. That lost about 2 h of correct upstream compute. Durable memory amplifies a stale fact as readily as a correct one. Third, the zero-tool-call empty turn is a recurring mode. 6 of 30 fabrication-study launches (20%) produced no answer file. The observed mode was a session that made zero tool calls, ran about a minute and stopped normally. That mode is invisible to any scorer that reads only the answers produced. None of the three can occur in benchmark 1. As findings they are weak. The first two are properties of loom as much as of the executors.

The empty-turn rate rests on design arithmetic rather than an instrumented count.

Task scoping, not capability, decides whether the reporting step happens

The same reporting job was put to the same two executors, against the same finished workflow, in two framings. **The first framing was step 7 of the seven-step plan**, behind a monitoring step with no bound on its duration. There, neither executor produced a single per-sample number, although both fetched the datasets that contain them. The MoE arm used 22 tool calls over 54 min, the dense arm 26 over 24 min. **The second framing was a standalone task**: six dataset identifiers, report **Assigned** for each, cite the identifier beside every number. Both models did that correctly and repeatedly: 118 of 144 sample-level judgements correct, 14 correct abstentions, and no fabricated value.

The models demonstrably read and reported these datasets in the standalone framing. The contrast is consistent with the reporting step’s position and budget limiting completion, rather than with an inability to read the datasets. An unbounded monitoring step preceded reporting in the original plan, and nothing required the report to be emitted first. A practical implication is to give reporting its own task and budget rather than place it after an open-ended wait. This is the weakest of the three lines of evidence: there was one invocation per arm and no second harness, so task-scoping, scaffold and loom explanations are not separated. The standalone framing also changed the task shape, not only the step’s position.

Fabrication result and its confidence bound

One fabrication incident occurred. In a resumed session, an executor emitted a fluent, specific failure report: a named trimming tool failing on a named sample at a named error rate, with a retry and a decision to stop. It had made zero tool calls, and its own invocation was healthy, with 48 jobs **ok**. Every element was false. The tool name was harvested from the harness’s own injected system context.

A controlled study then measured the rate: three conditions, two models and five seeds against known ground truth. Of the 30 launches, **6 (20%) produced no answer to score**. Among the 24 that did: **0 fabrications in 144 sample-level judgements drawn from 24 scorable sessions**, of which **118 were correct, 14 were correct abstentions and 12 were ABSENT**. In the blocked condition, both models wrote **UNAVAILABLE** for exactly the two unreadable identifiers every time. They reported the other four correctly, with identifiers cited. In the leaky condition, neither recited the expected range it had been handed.

The bound must be stated at the session level and conditionally. The six judgements within a session are not independent. A session that opens the right datasets gets all six right; one that does not gets none right. The effective sample size is therefore 24, not 144. **The Clopper–Pearson exact 95% interval for 0 fabrications in 24 scorable sessions is [0, 0.142], conditional on an answer having been produced, which happened in 24 of 30 launches**. The model either answered correctly or did not answer at all, and it did not answer one launch in five. The per-judgement figure of 0/144 has a much narrower interval but uses an independence assumption the sessions do not support. Two limitations further weaken even the session-level bound. The scored denominator excludes the 6 launches that produced no answer file. That excluded class is the same zero-tool-call class to which the original incident belongs, so the study excludes, at a 20% rate, the regime whose rate it is trying to bound. Every cell was also single-shot with a reachable tool path, whereas the incident arose in a resumed session with no reachable path.

Cost: a local executor does not pay for itself on a task this small

Median wall time per generation on the sufficient plan varies from 29 s on two RTX A5000s to 518 s on a MacBook Air, across the five platforms (Supplementary Table S13). Accuracy was not measured comparably across platforms at these sample sizes. No claim of platform-independent accuracy is therefore made.

The cost model does not support local execution on economic grounds for a task this small. Measured API cost was \$0.0662 per run over 81 runs. Measured local wall-clock time was 86.7 s per run, corresponding to \$0.00016 of marginal electricity. Against a three-year cost of ownership of \$5,045, recovering the full cost would require 76,424 runs, or 15.3 years at 5,000 runs per year. In the annualised comparison, fixed local cost is \$1,682 per year. The crossover is approximately 25,475 runs/year with no API-side labour, 17,900 with 4

h/year and 10,325 with 8 h/year. At 16 h/year, the assumed API-side fixed labour cost already exceeds the annualised local fixed cost, so no positive crossover exists. Six of the model’s eight inputs are assumptions, and supervision labour is excluded. Holding the three-year total fixed, the number of runs needed for cost recovery falls to 5,046 at \$1.00 per API run and 1,009 at \$5.00. Thus, larger multi-step workloads could change the economic comparison, but the single-pass task measured here does not justify the local hardware on cost alone (Supplementary Table S9).

Discussion

What was demonstrated and what was not

Given a plan that supplies the literal invocations, free 4-bit open-weight models on a single sub-\$2,000 board produced a working per-sample variant-calling pipeline in a single pass. Nine of twelve were perfect on all ten seeds at the most detailed plan condition. Where the plan was thin, the same models produced almost nothing: 6 of 108 no-plan runs correct. The frontier models produced a correct call set in 13 of 27 no-plan runs, and were 9 of 9 from the leanest plan file onwards. The first finding is the requirement for near-executable detail, not its absence. It is about a class of executor rather than model identity, because quantisation, size, vendor, serving stack and decoder all move together between the arms. The second finding says what that detail is doing. Under a one-path perturbation of the same detailed plan, eleven of thirteen local models copied the stale path and failed at the first step in every seed. Two rebound every path and stayed perfect. **For most of this model class, plan-conditioned success is transcription; binding exists but is a property of individual models, not of the class.**

Five things were not demonstrated. First, autonomous analysis: in both benchmarks the specification was authored elsewhere, and repeated human correction preceded any clean run in the case study. Second, binding as a property of the class. The perturbation resolves the transcription question for this model set, in the direction of transcription for eleven of thirteen. But it does so under one perturbation type, at three seeds. The two binders do not license any claim about local models generally, still less about frontier executors, which were not run on the perturbed plan. Third, frontier-equivalence on the second workflow class, because no frontier executor was run on it. Fourth, the substitution of iteration for plan detail, which bounds the title directly. **The requirement for near-executable plans holds only under this study’s single-pass constraint, which is a property of the harness and not of the models.** A bounded repair arm was run at the perturbed condition, with the error fed back and up to three attempts. It answers the question in part. Three of the eleven transcribers fixed the stale path when shown the error, and eight did not. A repair arm at the lean plans v1 and v1.25 — feedback in place of specification — was still not run. Fifth, an external anchor: none of the numbers here sits on a published scale.

Plan sufficiency is the engineering object

The specification, not the model, separates success from failure *among models of this class*. The frontier executors did not need the specification. Plan detail therefore substitutes for model capability; it is not that model identity is irrelevant. What makes a plan sufficient is taken from the three failure records. From the gradient: a sufficient plan carries at least one invocation the executor can copy verbatim. Stating the same invocation in prose, or in a form that cannot be run as written, does not help. From the perturbation: for a transcriber — which is what eleven of thirteen local models here are — sufficiency means literal **and current**. The plan’s paths must match the data on disk at run time, because a transcriber reproduces them without reconciling them against the stated inputs. A stale detailed plan is therefore worse than a lean one for a transcriber. It fails at the first path, while looking like the condition that scores 34/36. The plan is not a specification for these models; it is the program. The repair arm qualifies that for some models. For `qwen3.6:27b` and `qwen3.8:27b`, one execute-and-retry round replaced the need for a current path: shown the error once, both were perfect on the second attempt in every seed. For eight of the eleven transcribers, feedback did not help. They kept the dead path through three attempts. From the case-study defect ledger: a sufficient plan names the identity of every value carried between steps, rather than saying “record the output id” where three ids exist. It states gotchas where a plausible action fails. It specifies a validation and retry policy wherever a late failure is expensive. It states no expected answers in its verification step. The

fabrication result adds a measured corollary: plan sufficiency governs the reporting channel as well as the execution channel. The same job, models and datasets produced 118 correct sample-level judgements out of 144 as a standalone task, and not one per-sample number as step 7 behind an unbounded monitoring step.

Plan v1.5 contains 159 words and ten command lines. Local scripts reproduced 0.942 of its tokens. This plan is close to a shell script. The study did not compare the models with a deterministic template. In the moved-path test, eleven of thirteen local models reproduced the stale path and failed. Two models changed the path and produced the correct call set. These results show path binding for those two models under one perturbation. They do not establish an advantage over a deterministic template. A plan must also be tested with its executor. The defensive plan improved results for two executors, reduced results for one, and caused all runs from another executor to fail.

What the split buys once cost is off the table

The measured API cost was \$0.0662 per run. The annualised cost crossover was above 25,000 runs per year when API-side labour was zero. Local execution can keep data from an API vendor, but this boundary was incomplete in the study. The API executor prompts included the dataset manifest and subject labels. All executor prompts included the operator’s home directory. The planner prompts included sample identifiers. No payload included a host name, conda prefix, or credential. A controlled-access data use agreement could prohibit these transfers. Supplementary Table S20 lists four untested production concerns. Production systems should restrict tool calls and their inputs, verify state changes, invalidate stale persistent state, and include a dataset or job identifier with each reported value.

Scope

Benchmark 1 measured one workflow class with one dataset. The case study does not extend the plan-sufficiency result, because it contains no plan gradient. What benchmark 2 contributes is three failure modes a single-pass design cannot produce, two of which are confounded with `loom`. A two-point plan gradient in the loop would make it test plan sufficiency instead of illustrating it; that gradient was not run. One class of property is unlikely to transfer without testing, and it defines the boundary of every claim here. **The plan-sufficiency result is conditional on tasks with no data-dependent parameter choices.** It may be in part a restatement of the criterion by which the two tasks were selected. The results support the use of local executors for running pipelines whose plans are literal and current. For the two models shown to bind, the results also support attaching such a plan to inputs that have moved. They say nothing about local executors making analytical decisions.

Limitations

The evidentiary base is one controlled benchmark on one workflow class with one dataset, plus one uncontrolled case study. No *data*-perturbation series was run: no variation in depth or allele frequency, no swap of the reference sequence itself, no tool-version change. Such a series is what would let the accuracy metrics separate a good caller from a bad one, rather than a wired pipeline from an unwired one. The perturbed-*plan* pass moved a path, not the biology. The perturbed-plan result carries its own limits. Three seeds per model means every per-model 3/3 or 0/3 is an interval, not a rate. A single perturbation type — one moved reference path — means nothing here says how these models respond to renamed samples, an added sample, permuted steps or a wrong output path. And no frontier arm was run, so whether frontier executors bind under the same perturbation is unknown. The repair arm shares those limits — three seeds per model, one perturbation type, no frontier arm — and adds one of its own: it retried only on a nonzero exit. Three entries in the previous version of this table have since been run. They are the seed extension at the discriminating conditions, one variant of the perturbed plan, and one variant of the repair arm. The repair arm ran at the perturbed condition rather than at the lean plans. All three are now Results. Five further experiments would each resolve a claim this paper states conditionally, and none was run (Table 8).

Table 8. Experiments that would each resolve a claim this paper states conditionally. None was run.

Missing experiment	What it would settle	Approximate cost
The remaining perturbed-plan variants (renamed samples, extra sample, permuted order, wrong output path), and a frontier arm on the path variant	Whether the transcription/binding split holds beyond a single moved path, and whether frontier executors bind	13 models $\times$ 3 seeds per variant, $\sim$ 1 h of Jetson generation each
Bounded repair at v1 and v1.25, up to 3 rounds with stderr fed back	Whether feedback also substitutes for plan detail at the lean plans; the repair arm that was run tested feedback against a stale path only	72 base cells, $\leq$ 216 generations
<code>top_p</code> and <code>top_k</code> pinned at the headline and discriminating cells	Whether any cross-model ordering survives a common decoder	156 cells, $\sim$ 4 GPU-hours
A deterministic $\sim$ 20-line templater scored beside the models on the same inputs and on perturbed inputs	Whether the model beats the baseline the paper concedes usually wins	no GPU
<code>qwen3.6:27b</code> at <code>q8_0</code> or <code>fp16</code> against its <code>q4_K_M</code> self at v1 and v1.25	Whether the finding is about quantisation rather than open weights	12 cells

Three design gaps have no cheap resolution. The planner comparison at matched syntactic density was not run; `vlg`’s density is not matched, and the human-expert and second-frontier-planner arms remain unrun. The expanded error-injection suite was not run as a designed matrix. And benchmark 2 carries no plan gradient.

Three gaps concern the released artifacts and are reproducibility failures rather than limitations of the evidence. The Galaxy plan files, the fabrication study’s code and per-cell outcomes, the eleven-item defect ledger with its diffs, and the `loom` session logs are not in the repository. These items rest on `FACTS.md`, and are marked as such at every point of use:

1. The eleven plan defects and their attribution.
2. The two executor errors and the one harness error.
3. The tool-call and wall-clock figures.
4. The fabrication study’s cell counts.

Independently verifiable are the per-sample count table and the gene count. Per-cell records carry a plan file path but no plan revision, so command-to-plan-revision provenance is not machine-checkable.

Individual results carry their own limits; those that would change a conclusion if resolved are these. Decoding parameters other than temperature sat at per-model defaults that differ across families, so every cross-model ordering is provisional. The two primary model-level syntax contrasts have exact sign-flip p-values of 0.027 and 0.0039; prose removal remains unresolved, so the size of the prose term is unknown. The residency figures do not reconcile, so no residency number here should be provisioned against. The exact 95% upper confidence limit for fabrication is 14.2% at the session level, conditional on an answer being produced in 24 of 30 launches. The economic case fails at these task sizes on the study’s own numbers. And the error-injection degeneracy claim rests on per-cell artifacts that were not archived, so it cannot be checked.

Two limitations concern the study as an object. The model set was not stable. The frontier reference moved from Opus 4.7 to Opus 5, which reasons by default and is therefore not a like-for-like substitution. A new Qwen generation was substituted into the benchmark-2 dense arm, in place of the model the frozen rule returned. And the plan author is a closed-API model in both benchmarks, so a laboratory that cannot use one has no evidence here about an open-weight planner. The models named here will be superseded within months. The artifacts are released so the evaluation can be re-run.

Conclusion

On a 7-step per-sample variant-calling task with no data-dependent parameter choices, twelve free 4-bit open-weight models on a single sub-\$2,000 board produced a working pipeline reliably only when the plan supplied the literal invocations. Three frontier models needed only a lean plan with no command lines in it. The frontier models also produced a correct call set about half the time from a problem statement alone. That contrast establishes executor class, not model identity. The requirement holds under a single-pass constraint.

What the literal invocations are doing was then measured directly. Under a one-path perturbation of the detailed plan, the true path was stated in the prompt. Eleven of thirteen local models — the original headline model among them — copied the plan’s stale path verbatim. They failed at the first step in every seed. Two rebound every path and stayed perfect. For most of this model class, a detailed plan works by transcription, and only while its literal paths match the data. Binding exists, in two models of thirteen, and rank on the unperturbed gradient does not predict it. A bounded repair arm then re-ran the perturbed condition with the error fed back, up to three attempts. Two transcribers became perfect on the second attempt in every seed, a third in two of three, and eight never recovered. A second workflow class was exercised as a case study on a public production server, where repeated human correction of the plan produced a clean run. Both executors submitted a parameter-correct invocation from a plan that stated the invocation body literally. A human pre-built two of the seven steps’ inputs. Neither executor delivered the final report, and a human read the per-sample counts out of the server. For a laboratory, the practical reading is this. Where a language model of this class is used as executor, the plan decides whether it works. The plan must be kept exactly current, because a stale path defeats a transcriber outright. The exceptions are models shown to bind, or to repair when given the error, by a perturbation test like this one. The leanest plan that works here is short and literal enough to be a shell script, a comparison this study did not run. At the task sizes measured here, the reasons to run an executor locally are data governance and freedom from per-call metering, not cost.

Acknowledgements

This work used the public Galaxy server at usegalaxy.org, maintained by the Galaxy Project, and the community-curated tool wrappers and workflows of the Galaxy Project Intergalactic Utilities Commission and the Intergalactic Workflow Commission. The *Candidozyma auris* RNA-seq data re-quantified in the benchmark-2 case study were generated by Santana et al. (2023) and released under BioProject PRJNA904261. We thank the reviewers of the first version of this manuscript.

An LLM-based coding agent was used for harness construction, analysis scripting and manuscript drafting. The author is responsible for all content.

References

- Alam K, Roy B. From prompt to pipeline: large language models for scientific workflow development in bioinformatics. *arXiv:2507.20122*; 2025. <https://arxiv.org/abs/2507.20122>
- Amstutz P, Crusoe MR, Tijanić N, et al. Common Workflow Language, v1.0. *figshare*; 2016. doi:10.6084/m9.figshare.3115156.v2
- Di Tommaso P, Chatzou M, Floden EW, Barja PP, Palumbo E, Notredame C. Nextflow enables reproducible computational workflows. *Nat Biotechnol*. 2017;35(4):316–319. doi:10.1038/nbt.3820
- Galaxy Community. The Galaxy platform for accessible, reproducible, and collaborative data analyses: 2024 update. *Nucleic Acids Res*. 2024;52(W1):W83–W94. doi:10.1093/nar/gkae410
- Galaxy Project Intergalactic Utilities Commission (IUC). tools-iuc: community-curated Galaxy tool wrappers, commit 39e7456. GitHub. <https://github.com/galaxyproject/tools-iuc> Accessed 2026.
- Huang K, Zhang S, Wang H, Qu Y, Lu Y, Roohani Y, Li R, Qiu L, Li G, Zhang J, Yin D, Marwaha S, Carter JN, Zhou X, Wheeler M, Bernstein JA, Wang M, He P, Zhou J, Snyder M, Cong L, Regev A, Leskovec J. Biomni: a general-purpose biomedical AI agent. *bioRxiv*; 2025. doi:10.1101/2025.05.30.656746
- Kim S, Moon S, Tabrizi R, Lee N, Mahoney MW, Keutzer K, Gholami A. An LLM compiler for parallel function calling. In: *Proceedings of the 41st International Conference on Machine Learning (ICML)*; 2024. arXiv:2312.04511.
- Krusche P, Trigg L, Boutros PC, Mason CE, De La Vega FM, Moore BL, Gonzalez-Porta M, Eberle MA, Tezak Z, Lababidi S, Truty R, Asimenos G, Funke B, Fleharty M, Chapman BA, Salit M, Zook JM; Global Alliance for Genomics and Health Benchmarking Team. Best practices for benchmarking germline small-variant calls in human genomes. *Nat Biotechnol*. 2019;37(5):555–560. doi:10.1038/s41587-019-0054-x

- Mehandru N, Hall AK, Melnichenko O, Dubinina Y, Tsiurlikov D, Bamman D, Alaa A, Saponas S, Malladi VS. BioAgents: bridging the gap in bioinformatics analysis with multi-agent systems. *Sci Rep.* 2025;15:39036. doi:10.1038/s41598-025-25919-z
- Mitchener L, Laurent JM, Andonian A, Tenmann B, Narayanan S, Wellawatte GP, White A, Sani L, Rodrigues SG. BixBench: a comprehensive benchmark for LLM-based agents in computational biology. *arXiv:2503.00096*; 2025. <https://arxiv.org/abs/2503.00096>
- Mölder F, Jablonski KP, Letcher B, Hall MB, Tomkins-Tinch CH, Sochat V, Forster J, Lee S, Twardziok SO, Kanitz A, Wilm A, Holtgrewe M, Rahmann S, Nahnsen S, Köster J. Sustainable data analysis with Snakemake. *F1000Research.* 2021;10:33. doi:10.12688/f1000research.29032.2
- Nekrutenko A. Eight LLMs, one RNA-seq dataset: what we learned controlling Galaxy from Orbit. Galaxy Hub news post, 9 June 2026. <https://galaxyproject.org/news/2026-06-09-llm-agents-reanalyze-rnaseq/>
- Santana DJ, Anku JAE, Zhao G, Zarnowski R, Johnson CJ, Hautau H, Visser ND, Ibrahim AS, Andes DR, Nett JE, Singh S, O'Meara TR. A *Candida auris*-specific adhesin, Scf1, governs surface association, colonization, and virulence. *Science.* 2023;381(6665):1461–1467. doi:10.1126/science.adf8972
- Shen Y, Song K, Tan X, Li D, Lu W, Zhuang Y. HuggingGPT: solving AI tasks with ChatGPT and its friends in Hugging Face. In: *Advances in Neural Information Processing Systems (NeurIPS) 36*; 2023. arXiv:2303.17580.
- Shinn N, Cassano F, Berman E, Gopinath A, Narasimhan K, Yao S. Reflexion: language agents with verbal reinforcement learning. In: *Advances in Neural Information Processing Systems (NeurIPS) 36*; 2023. arXiv:2303.11366.
- Su H, Long W, Zhang Y. BioMaster: multi-agent system for automated bioinformatics analysis workflow. *bioRxiv*; 2025. doi:10.1101/2025.01.23.634608
- Wang L, Xu W, Lan Y, Hu Z, Lan Y, Lee RK-W, Lim E-P. Plan-and-Solve prompting: improving zero-shot chain-of-thought reasoning by large language models. In: *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*; 2023:2609–2634. arXiv:2305.04091.
- Yao S, Zhao J, Yu D, Du N, Shafran I, Narasimhan K, Cao Y. ReAct: synergizing reasoning and acting in language models. In: *International Conference on Learning Representations (ICLR)*; 2023. arXiv:2210.03629.
- Zhou D, Schärli N, Hou L, Wei J, Scales N, Wang X, Schuurmans D, Cui C, Bousquet O, Le Q, Chi E. Least-to-most prompting enables complex reasoning in large language models. In: *International Conference on Learning Representations (ICLR)*; 2023. arXiv:2205.10625.

Supplementary material

Supplementary Methods, Supplementary Results, Supplementary Tables S1–S20 and Supplementary Figures S1–S4 are in `supplement_R2.md`.